

Included Re-ranking – rerankscore is seen here

The screenshot displays the Retrieval application interface and its network logs. The interface features a search bar labeled 'Retrieval' with the subtitle 'Search across every modality at once'. Below the search bar are filters for 'Text', 'Image', 'Multilingual', 'Audio', and 'Video'. A 'Text query' section includes a 'RAG' input field and a 'Search' button. The 'Query usage' section indicates 'Prompt: 5 - Total: 5'. Three search results are shown, each with a 'TEXT' label, a score, and a summary. The first result has a score of 0.787, the second 0.716, and the third 0.705. The network logs on the right show the API request and response, including the 'search' endpoint and the JSON response containing document metadata and content.

Included De-duplication:

The screenshot displays the Testleaf application running on a local host (localhost:3000). The interface is divided into several sections:

- Testleaf Multimodal RAG:** A green banner stating "Ingest and search across one shared vector space — or generate new text, image, and audio content instantly, no ingestion required." Below this are buttons for Text, Image, Multilingual, Audio, and Video.
- Instant AI Summaries:** A section with a purple lightning bolt icon, stating "Every retrieved result can be summarized on demand via Groq — turning a raw match into a clear, readable answer in seconds."
- Ingestion:** A section titled "Add content to the shared vector store" with buttons for Text, Image, Multilingual, Audio, and Video. Below these is a text input field labeled "Text" with the placeholder "Enter text to ingest...". A character count shows "0 characters · ~1 tokens (estimate)". An "Ingest Text" button is present.
- Notification:** A green box at the bottom states: "Ingested successfully — id: 6a93c4def4e98be023875e1a, category: text. Note: identical content already exists (id: 6a93c4def4e98be023875e1a)."

The browser's developer tools (Network tab) show a successful POST request to the ingestion endpoint. The response is a JSON object:

```
{
  "data": {
    "id": "6a93c4def4e98be023875e1a",
    "category": "text",
    "duplicateOf": "6a93c4def4e98be023875e1a",
    "duplicate": true,
    "dimensions": 1824,
    "provider": "jina",
    "model": "jina-embeddings-v5-omni-small",
    "costUsd": 0
  },
  "message": "Duplicate item skipped. No new emb"
}
```


Generation – Based on selected service providers with fallbacks added

The screenshot shows a web application running on localhost:3000. The interface has tabs for Text, Audio, Video, and Image. The Text tab is active, showing a text input with "Hello Good morning" and a "Generate" button. Below the input, there are two output sections: "Image" (labeled CLOUDFLARE) showing a generated image of a person on a cliff, and "Audio" (labeled GROQ) showing a play button and a progress bar. To the right, the browser's network tab is open, showing a request to "generate" with a response containing JSON data for both image and audio outputs.

```
1 {
  "success": true,
  "data": {
    "inputModality": "text",
    "outputs": {
      "image": {
        "status": "done",
        "resultDataUrl": "data:image/png;base64...",
        "provider": "cloudflare",
        "model": "cf/black-forest-labs/7",
        "costUsd": 0
      },
      "audio": {
        "status": "done",
        "resultDataUrl": "data:audio/mp3;base64...",
        "provider": "groq",
        "model": "canopy",
        "costUsd": 0
      }
    }
  },
  "message": "Generation complete."
}
```

The screenshot shows the same web application with the Audio tab active. The "INPUT METHOD" section has "Upload File" selected, and a file named "Meeting Recording.mp3" is uploaded. Below the file, there is a "Generate from Audio" button. The "Text" output section shows a generated text transcript of the audio. The "Image" output section (labeled CLOUDFLARE) shows a generated image with the text "DAILY STAND-UP" and a list of bullet points. To the right, the browser's network tab is open, showing a request to "generate" with a response containing JSON data for both text and image outputs.

```
1 {
  "success": true,
  "data": {
    "inputModality": "audio",
    "outputs": {
      "text": {
        "status": "done",
        "text": "Hi, everyone. I'm Hari. I'm going to run the daily stand-up meeting today. We can get started. Morning, Hari. Morning, Sindhu. Um, Radhakrishnan here. So, yesterday I've been assigned with a couple of tickets related to balance transfer. So I have written manual test cases initially for that and got it signed off from the BA and PO in our team. And since the latest build is now available in QA, I've started testing those scenarios and monitoring how the system behaves for the different scenarios. So far, the system is good and there are no any defects and if there are any, I will raise it accordingly. And once the execution is done, I will focus on converting those steps into automation scripts and see how it goes. So that's the update from my side. Thanks. Thank you, Radhakrishnan. Good morning, everyone. Sindhu here. Yesterday I've completed smoke and regression testing for the payments module and logged three defects in Jira. And today I will be continuing manual testing for the reports module and automating the transaction history flow in Selenium. And... and no... and no blockers so far and waiting for a new build to proceed with re-validation. So that's all for my update. Thank you. Okay, thank you, Sindhu. Thank you, guys. Thanks for your update. Thanks, everyone. Thank you."
      },
      "image": {
        "status": "done",
        "resultDataUrl": "data:image/png;base64...",
        "provider": "cloudflare",
        "model": "cf/black-forest-labs/7",
        "costUsd": 0
      }
    }
  },
  "message": "Generation complete."
}
```



Token details added:

TEXTscore: 0.705

Testing the multimodal RAG system with provider and model metadata preservation through the reporting layer.

lang: en · encoding: UTF-8 · source: text-input

Re-summarize ▾ Hide summary

The study evaluates a multimodal retrieval-augmented generation (RAG) system. It focuses on preserving provider and model metadata, using a reporting layer to monitor and verify that the metadata remains intact throughout the process.

Summary tokens — Total: 294 · Input: 78 · Output: 216 · Reasoning: 170

User can select the required service providers and embedding models. Also shows the model details:

**Testleaf**
ALWAYS AHEAD

PROVIDER **Cohere** EMBEDDING **Jina** 

**Generate**
Create the other modalities from one input — works standalone, no vector store needed

VOICE **Gemini** IMAGE **Cloudflare**

Text Audio Video Image

Text
help!!

Generate

Image
MODEL: @CF/BLACK-FOREST-LABS/FLUX-1-SCHNELL · PROVIDER: CLOUDFLARE · COST: \$0.000000



Audio
MODEL: GEMINI-2.5-FLASH-PREVIEW-TTS · PROVIDER: GEMINI · COST: \$0.000000

